

# Project 1: Emotion-Change Detection for Classroom Feedback

*Initial Proposal • October 27, 2025*

**To-Do:** Add the specific Kaggle dataset URL and citation once selected.

## 1) Motivation

**(A) Problem.** In typical learning environments, instructors lack accurate, real-time feedback about how their teaching methods affect students' emotions. It's nearly impossible to track class-wide emotional responses as they happen and do so reliably.

**(B) Why this matters for AI/ML.** Modern ML methods can quantify qualitative cues and reveal hidden patterns in emotional responses to instruction. Turning facial expressions into measurable signals enables evidence-based adjustments to teaching.

**(C) Real-world motivation.** Personalized teaching improves when educators can see which methods foster positive, sustained engagement. Better feedback loops can support equity by helping instructors adapt techniques to individual needs.

## 2) Task Definition

### (A) Inputs → Outputs.

- **Input:** two facial images of the *same student* captured at different times (e.g., snapshots every ~10 s).
- **Output:** JSON with the change (delta) in confidence for each emotion: `anger`, `disgust`, `fear`, `happiness`, `neutral`, `sadness`, `surprise`.

### (B) Concrete I/O example (Python-style pseudocode).

```
# Not executable; senior-engineer-readable pseudocode
def detect_emotion_change(image_path_1: str, image_path_2: str) -> dict:
    # 1) Load images from paths (same student, different times)
    # 2) Run both through a pre-trained facial emotion classifier
    #     -> returns per-emotion confidence scores for each image
    # 3) Compute deltas: scores_t2 - scores_t1 for each of 7 emotions
    # 4) Return JSON-like dict of deltas (no persistence, no PII)

    return { # example shape
        "anger": d_anger,
        "disgust": d_disgust,
        "fear": d_fear,
        "happiness": d_happiness,
```

```

    "neutral": d_neutral,
    "sadness": d_sadness,
    "surprise": d_surprise
}

```

**(C) Task type.** Multi-class **classification** (emotion recognition), summarized as per-class deltas between two snapshots.

### 3) Baseline

Use a well-tested, off-the-shelf facial emotion classifier to score each image separately, then compute per-emotion deltas ( $t_2 - t_1$ ). Baseline ends at delta computation; integration with lecture timelines is **out of scope** for Project 1.

### 4) Proposed Approach

- Adopt a vetted foundation model for facial emotion recognition (e.g., transformer-based) with documented evaluations.
- Record why it is appropriate (reported accuracy, training data traits, published benchmarks); cite model card/paper.
- Implement a minimal inference wrapper that: (i) validates two input paths, (ii) runs emotion scoring, (iii) returns JSON deltas.
- *(Optional stretch)* Accept video input, auto-sample frames, and alert only on large deltas. (Not required for Project 1.)

### 5) Evaluation Plan

**(A) Dataset.** *TBD Kaggle dataset (facial emotion classification).* Add URL/citation once finalized.

#### (B) Metrics.

- **Macro-F1 (primary):** average F1 across the seven emotions (per-image classification).
- **Delta accuracy (secondary):** where paired ground-truth labels exist, report mean absolute error between predicted and true deltas per class.
- **Note:** ROC-AUC is less informative here due to the multi-class setup and is not a primary metric.

#### (C) Comparisons.

- 1) Report benchmark metrics from the model card on the chosen dataset.
- 2) External sanity check: a black-and-white facial-expression poster set (e.g., “Ernest/Varney” style), resized/normalized as a small out-of-distribution probe.

#### (D) Quantitative & qualitative notes.

- *Quantitative:* per-class confidence scores and macro-F1.
- *Qualitative limitations:* labeler bias; single-label simplification (mixed emotions not modeled); cultural/context variability in expressions.

## 6) Plan & Milestones (week-by-week sketch)

- **Week 1:** Select model and dataset; set up inference; reproduce minimal baseline on a small batch.
  - **Week 2:** Implement delta computation and JSON output; add basic input validation/logging.
  - **Week 3:** Run evaluations (macro-F1, delta checks); document limitations and error cases.
  - **Week 4:** Polish write-up (figures/tables), package code repo, and finalize citations.
- 

## Appendix A: Brief Metrics Primer

- **Precision:** of predicted positives for a class, the fraction that are correct.
- **Recall:** of actual positives for a class, the fraction correctly identified.
- **F1:** harmonic mean of precision and recall; **macro-F1** averages F1 equally across classes.
- **Why not ROC-AUC?** Most natural for binary classification; multi-class ROC-AUC (one-vs-rest) is less informative here than macro-F1.
- **Delta accuracy:** summarize mean/median absolute error of predicted vs. true deltas across paired samples.